

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df = pd.read_csv(r"C:\Users\VICTUS\Desktop\Food_Delivery_Times (1).csv")
```

```
In [3]: df.head()
```

```
Out[3]:
```

	Order_ID	Distance_km	Weather	Traffic_Level	Time_of_Day	Vehicle_Type	Preparation_Time_min	Courier_Experience_yrs	Delivery_Time
0	522	7.93	Windy	Low	Afternoon	Scooter	12	1.0	
1	738	16.42	Clear	Medium	Evening	Bike	20	2.0	
2	741	9.52	Foggy	Low	Night	Scooter	28	1.0	
3	661	7.44	Rainy	Medium	Afternoon	Scooter	5	1.0	
4	412	19.03	Clear	Low	Morning	Bike	16	5.0	



```
In [4]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Order_ID              1000 non-null  int64
1   Distance_km           1000 non-null  float64
2   Weather               970 non-null   object
3   Traffic_Level         970 non-null   object
4   Time_of_Day           970 non-null   object
5   Vehicle_Type          1000 non-null   object
6   Preparation_Time_min  1000 non-null  int64
7   Courier_Experience_yrs 970 non-null   float64
8   Delivery_Time_min     1000 non-null  int64
dtypes: float64(2), int64(3), object(4)
memory usage: 70.4+ KB

```

In [5]: `df.describe()`

```

Out[5]:

```

	Order_ID	Distance_km	Preparation_Time_min	Courier_Experience_yrs	Delivery_Time_min
count	1000.000000	1000.000000	1000.000000	970.000000	1000.000000
mean	500.500000	10.059970	16.982000	4.579381	56.732000
std	288.819436	5.696656	7.204553	2.914394	22.070915
min	1.000000	0.590000	5.000000	0.000000	8.000000
25%	250.750000	5.105000	11.000000	2.000000	41.000000
50%	500.500000	10.190000	17.000000	5.000000	55.500000
75%	750.250000	15.017500	23.000000	7.000000	71.000000
max	1000.000000	19.990000	29.000000	9.000000	153.000000

In [6]: `df.shape`

Out[6]: (1000, 9)

```
In [7]: df.duplicated()
```

```
Out[7]: 0      False
        1      False
        2      False
        3      False
        4      False
        ...
        995    False
        996    False
        997    False
        998    False
        999    False
        Length: 1000, dtype: bool
```

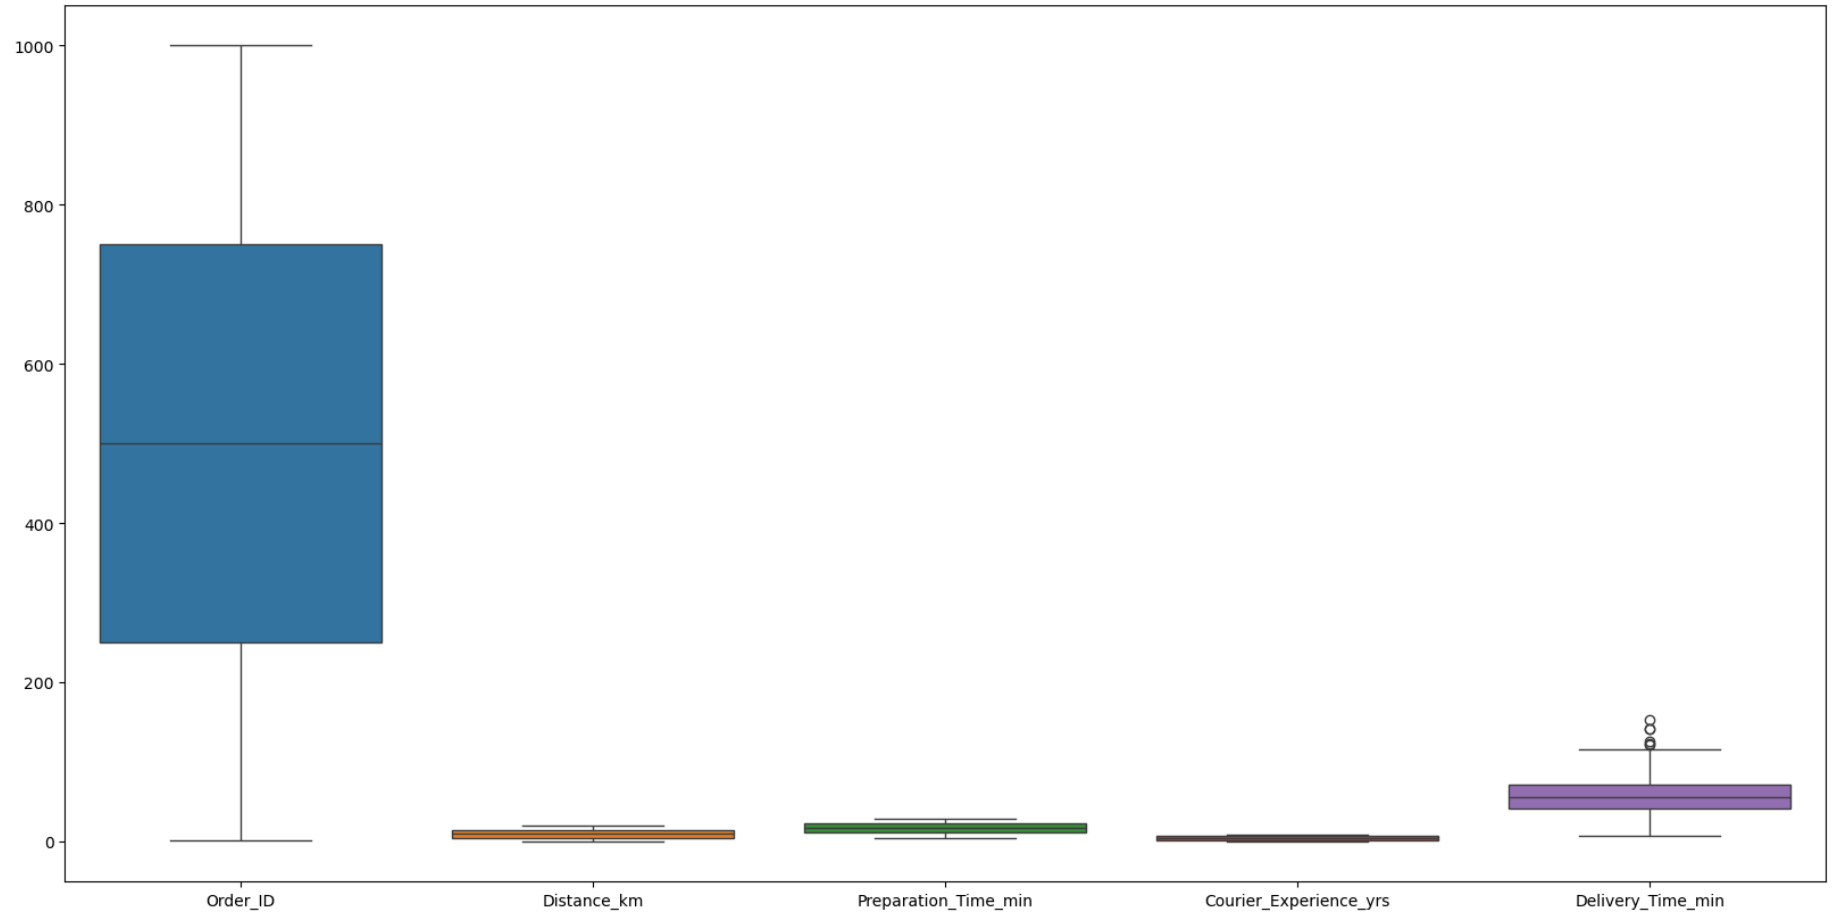
```
In [8]: df.isnull().sum()
```

```
Out[8]: Order_ID      0
        Distance_km    0
        Weather      30
        Traffic_Level  30
        Time_of_Day    30
        Vehicle_Type    0
        Preparation_Time_min  0
        Courier_Experience_yrs  30
        Delivery_Time_min  0
        dtype: int64
```

```
In [9]: def outliertreat(df,col):
        Q1 = df[col].quantile(0.25)
        Q3 = df[col].quantile(0.75)
        IQR = Q3 - Q1
        UL = Q3 + IQR
        LL = Q1 - IQR
        df.loc[df[col]>UL,col] = df[col].median()
        df.loc[df[col]<LL,col] = df[col].median()
```

```
In [10]: plt.figure(figsize=(20,10))
        sns.boxplot(df)
```

```
plt.show()
```



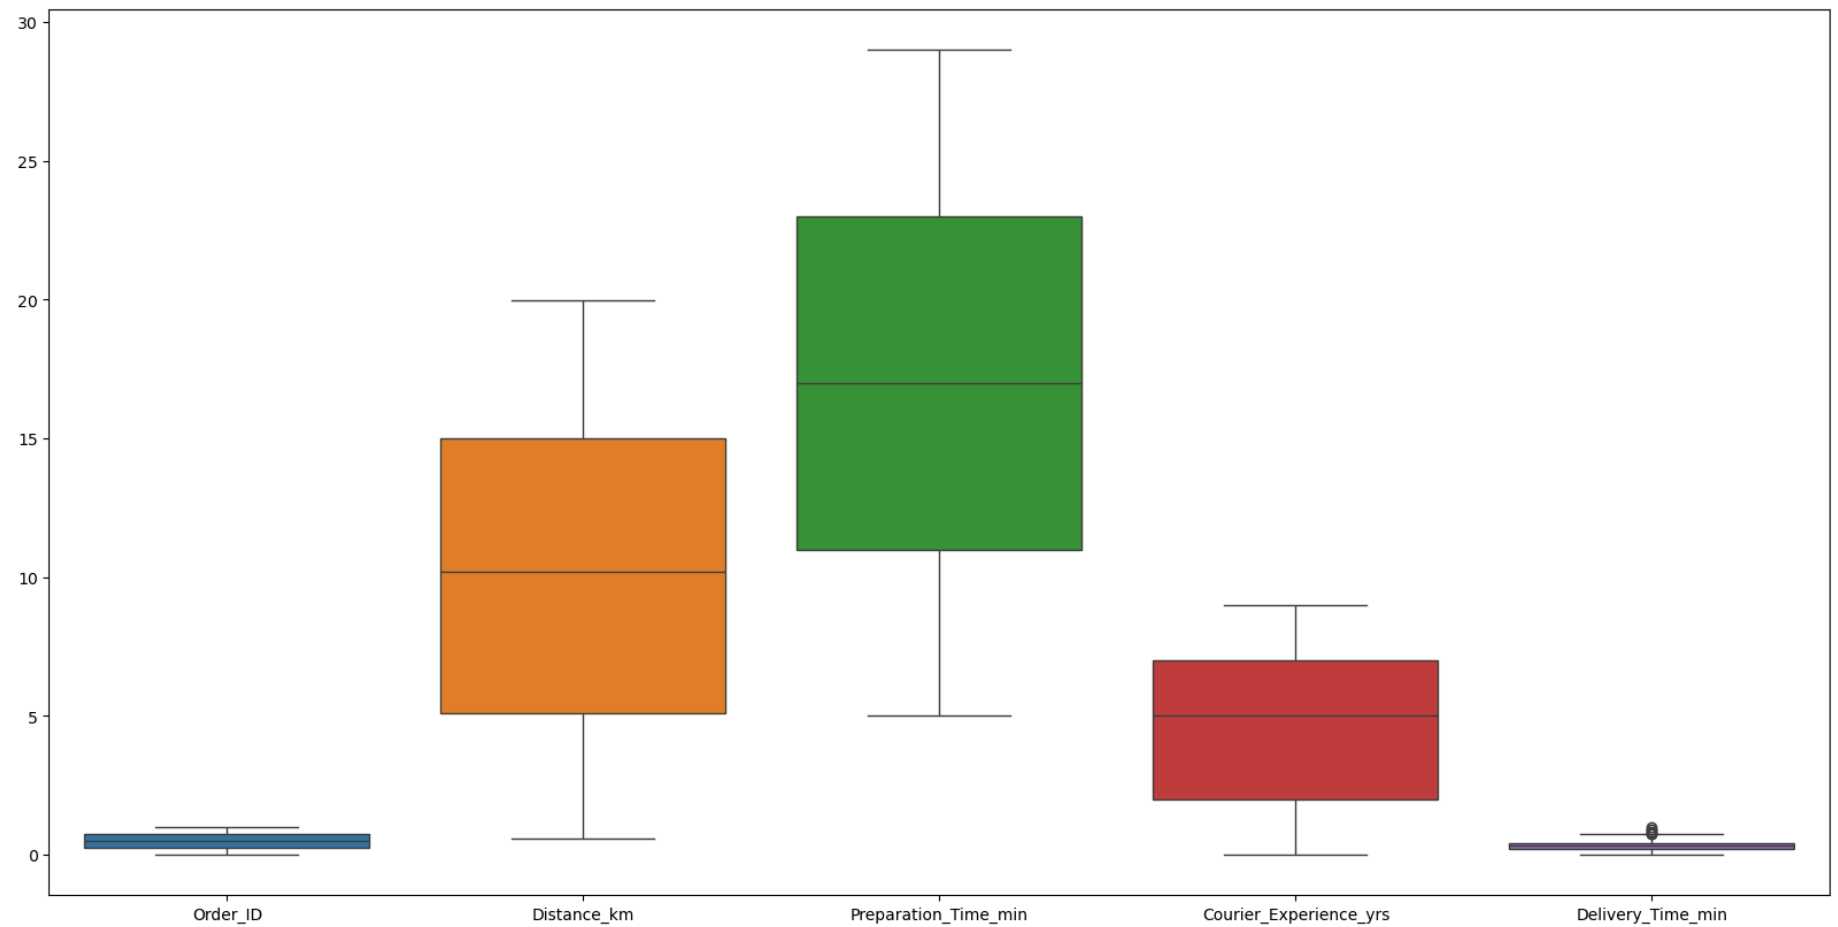
```
In [11]: from sklearn.preprocessing import MinMaxScaler
```

```
In [12]: MS = MinMaxScaler()
```

```
In [13]: df['Order_ID'] = MS.fit_transform(df[['Order_ID']])
```

```
In [14]: df['Delivery_Time_min'] = MS.fit_transform(df[['Delivery_Time_min']])
```

```
In [15]: plt.figure(figsize=(20,10))
sns.boxplot(df)
plt.show()
```



```
In [16]: from sklearn.preprocessing import LabelEncoder
```

```
In [17]: LE = LabelEncoder()
```

```
In [18]: df['Weather'] = LE.fit_transform(df['Weather'])
```

```
In [19]: df['Traffic_Level'] = LE.fit_transform(df['Traffic_Level'])
```


```
In [20]: df['Time_of_Day'] = LE.fit_transform(df['Time_of_Day'])
```

```
In [21]: df['Vehicle_Type'] = LE.fit_transform(df['Vehicle_Type'])
```

```
In [22]: df.head()
```

```
Out[22]:
```

	Order_ID	Distance_km	Weather	Traffic_Level	Time_of_Day	Vehicle_Type	Preparation_Time_min	Courier_Experience_yrs	Delivery_Time
0	0.521522	7.93	4	1	0	2	12	1.0	0.24
1	0.737738	16.42	0	2	1	0	20	2.0	0.52
2	0.740741	9.52	1	1	3	2	28	1.0	0.35
3	0.660661	7.44	2	2	0	2	5	1.0	0.20
4	0.411411	19.03	0	1	2	0	16	5.0	0.41

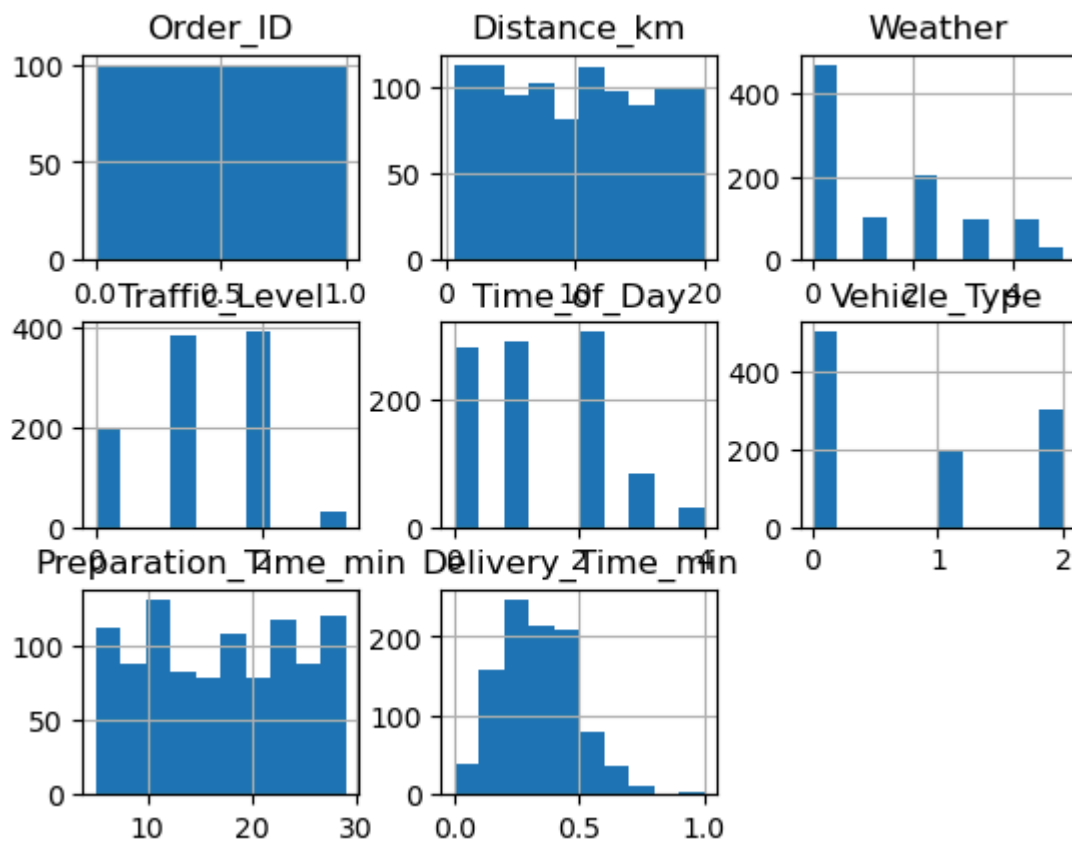


```
In [23]: df.drop(columns=['Courier_Experience_yrs'], inplace=True)
```

```
In [24]: df.isnull().sum()
```

```
Out[24]: Order_ID          0
Distance_km          0
Weather             0
Traffic_Level       0
Time_of_Day         0
Vehicle_Type        0
Preparation_Time_min 0
Delivery_Time_min    0
dtype: int64
```

```
In [25]: df.hist()
plt.show()
```



```
In [26]: df.skew(numeric_only=True)
```

```
Out[26]: Order_ID          -1.096752e-15
Distance_km      3.884047e-02
Weather          7.708209e-01
Traffic_Level    -1.383483e-01
Time_of_Day      4.518274e-01
Vehicle_Type     4.017149e-01
Preparation_Time_min  3.000816e-02
Delivery_Time_min  5.072512e-01
dtype: float64
```

```
In [27]: from sklearn.preprocessing import PowerTransformer
```

```
In [28]: PT = PowerTransformer()
```

```
In [29]: PT.fit_transform(df[['Traffic_Level', 'Preparation_Time_min', 'Delivery_Time_min']])
```

```
Out[29]: array([[ -0.31438592, -0.64489576, -0.57624906],
               [  0.93145906,  0.46559741,  1.21564784],
               [-0.31438592,  1.44926602,  0.19712054],
               ...,
               [-1.56398662,  1.21220682,  1.10374233],
               [-0.31438592, -1.27912648,  0.01385886],
               [-0.31438592,  0.96965152,  0.15189201]])
```

```
In [30]: df[['Traffic_Level', 'Preparation_Time_min', 'Delivery_Time_min', ]] = PT.fit_transform(df[['Traffic_Level', 'Preparation_Time_min', 'Delivery_Time_min']])
```

```
In [31]: df.skew(numeric_only=True)
```

```
Out[31]: Order_ID          -1.096752e-15
         Distance_km       3.884047e-02
         Weather           7.708209e-01
         Traffic_Level     -1.419601e-01
         Time_of_Day        4.518274e-01
         Vehicle_Type       4.017149e-01
         Preparation_Time_min -1.285052e-01
         Delivery_Time_min   8.900913e-03
         dtype: float64
```

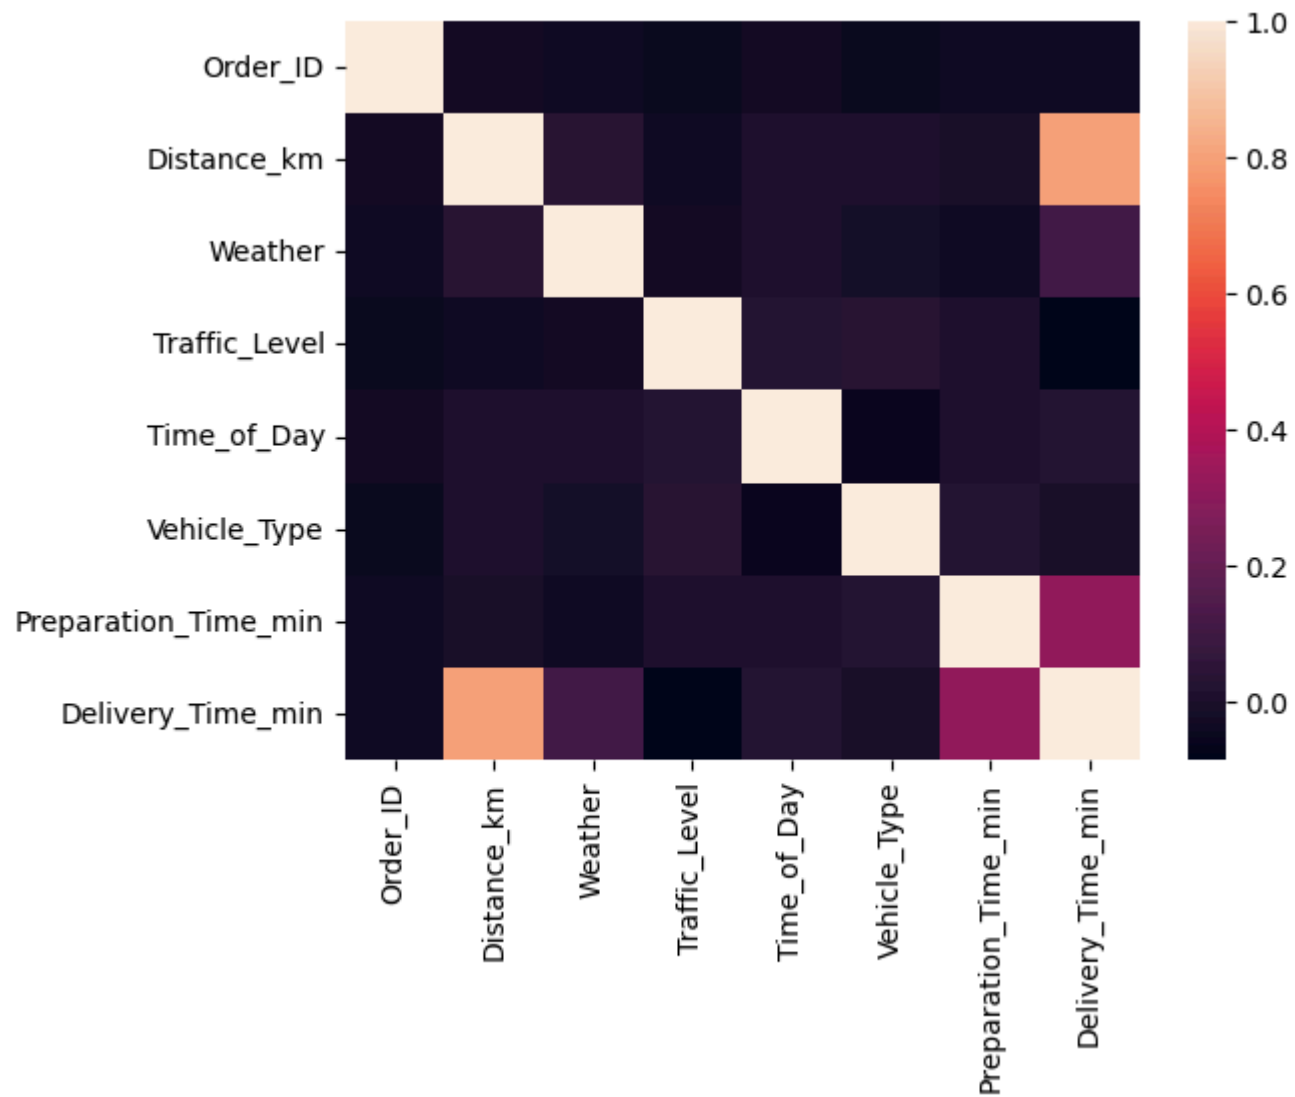
```
In [32]: df.corr(numeric_only=True)
```


Out[32]:

	Order_ID	Distance_km	Weather	Traffic_Level	Time_of_Day	Vehicle_Type	Preparation_Time_min	Delivery_Time_r
Order_ID	1.000000	-0.024483	-0.035785	-0.050806	-0.027034	-0.045030	-0.035859	-0.0328
Distance_km	-0.024483	1.000000	0.029756	-0.036625	0.009034	0.003319	-0.012127	0.7984
Weather	-0.035785	0.029756	1.000000	-0.031329	0.006595	-0.019231	-0.037662	0.1097
Traffic_Level	-0.050806	-0.036625	-0.031329	1.000000	0.022548	0.032533	0.004215	-0.0866
Time_of_Day	-0.027034	0.009034	0.006595	0.022548	1.000000	-0.054988	0.004447	0.0164
Vehicle_Type	-0.045030	0.003319	-0.019231	0.032533	-0.054988	1.000000	0.018661	-0.0042
Preparation_Time_min	-0.035859	-0.012127	-0.037662	0.004215	0.004447	0.018661	1.000000	0.3190
Delivery_Time_min	-0.032890	0.798499	0.109796	-0.086679	0.016430	-0.004227	0.319076	1.0000

```
In [33]: sns.heatmap(df.corr(numeric_only=True))
```

Out[33]: <Axes: >



```
In [34]: from sklearn.model_selection import train_test_split
```

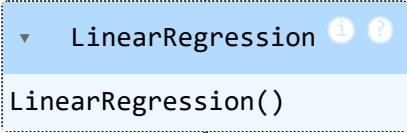
```
In [35]: X = df.drop('Delivery_Time_min',axis=1)
y = df.Delivery_Time_min
```

```
In [36]: X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.25,random_state=42)
```

```
In [37]: from sklearn.linear_model import LinearRegression
```

```
In [38]: LR = LinearRegression()
```

```
In [39]: LR.fit(X_train, y_train)
```

```
Out[39]: 
LinearRegression()
```

```
In [40]: LR_pred = LR.predict(X_test)
```

```
In [41]: from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

```
In [42]: mean_absolute_error(y_test,LR_pred)
```

```
Out[42]: 0.3354891852033775
```

```
In [43]: mean_squared_error(y_test,LR_pred)
```

```
Out[43]: 0.21446556795516905
```

```
In [44]: r2_score(y_test, LR_pred)
```

```
Out[44]: 0.793097088496856
```

```
In [45]: LR.score(X_train,y_train)*100
```

```
Out[45]: 74.54719561652814
```

```
In [46]: LR.score(X_test,y_test)+
```

```
Cell In[46], line 1
    LR.score(X_test,y_test)+
```

SyntaxError: invalid syntax

```
In [ ]: from sklearn.neighbors import KNeighborsRegressor
```

```
In [ ]: KNR = KNeighborsRegressor(n_neighbors=7)
```

```
In [ ]: KNR.fit(X_train,y_train)
```

```
In [ ]: KNR_pred = KNR.predict(X_test)
```

```
In [ ]: mean_absolute_error(y_test,KNR_pred)
```

```
In [ ]: mean_squared_error(y_test,LR_pred)
```

```
In [ ]: r2_score(y_test, KNR_pred)
```

```
In [ ]: KNR.score(X_train,y_train)*100
```

```
In [ ]: KNR.score(X_test,y_test)
```

```
In [ ]: from sklearn.tree import DecisionTreeRegressor
```

```
In [ ]: DTC = DecisionTreeRegressor(random_state=42)
```

```
In [ ]: DTC.fit(X_train , y_train)
```

```
In [ ]: y_pred_dt = DTC.predict(X_test)
```

```
In [ ]: mean_absolute_error(y_test, y_pred_dt)
```

```
In [ ]: mean_squared_error(y_test, y_pred_dt)
```

```
In [ ]: r2_score(y_test,y_pred_dt )
```

```
In [ ]: DTC.score(X_train,y_train)*100
```

```
In [ ]: DTC.score(X_test,y_test)
```

```
In [ ]:
```

```
In [ ]: from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
```

```
In [ ]: RFR=RandomForestRegressor(n_estimators=100, random_state=42)
```

```
In [ ]: RFR.fit(X_train, y_train)
```

```
In [ ]: RFR_pred = RFR.predict(X_test)
```

```
In [ ]: mean_absolute_error(y_test, RFR_pred)
```

```
In [ ]: mean_squared_error(y_test, RFR_pred)
```

```
In [ ]: r2_score(y_test, RFR_pred)
```

```
In [ ]: RFR.score(X_train,y_train)*100
```

```
In [ ]: RFR.score(X_test,y_test)
```

```
In [ ]: GBL = GradientBoostingRegressor(random_state=42, n_estimators=100, learning_rate=0.1)
```

```
In [ ]: GBL.fit(X_train, y_train)
```

```
In [ ]: gbl_pred = GBL.predict(X_test)
```

```
In [ ]: mean_squared_error(y_test, gbl_pred)
```

```
In [ ]: mean_absolute_error(y_test, gbl_pred)
```

```
In [ ]: r2_score(y_test, gbl_pred)
```

```
In [ ]: GBL.score(X_train, y_train)*100
```

```
In [ ]: GBL.score(X_test, y_test)
```

```
In [60]: from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
```

```
In [56]: param_grid = {  
        'n_estimators': [50, 100, 200],  
        'max_depth' : [None, 10, 20, 30],  
        }
```

```
In [58]: rf_model = RandomForestRegressor(random_state=42)
```

```
-----  
NameError                                Traceback (most recent call last)  
Cell In[58], line 1  
----> 1 rf_model = RandomForestRegressor(random_state=42)  
  
NameError: name 'RandomForestRegressor' is not defined
```

```
In [ ]:
```